

Anomaly Detection

5/21

Outline

- What is Anomaly Detection
- Classic Method
 - With Classifier
 - GMM (Gaussian Mixture Model)
 - Auto-Encoder
 - PCA
 - Isolation Forest
 - Summary
- Anomaly Detection on image
 - AnoGAN
 - EGBAD
 - GANomaly
 - Summary
- Anomaly Detection on Audio
 - GMGAN


What is Anomaly Detection

What is Anomaly

- Training Data



Anomaly



Anomaly

Classic Method

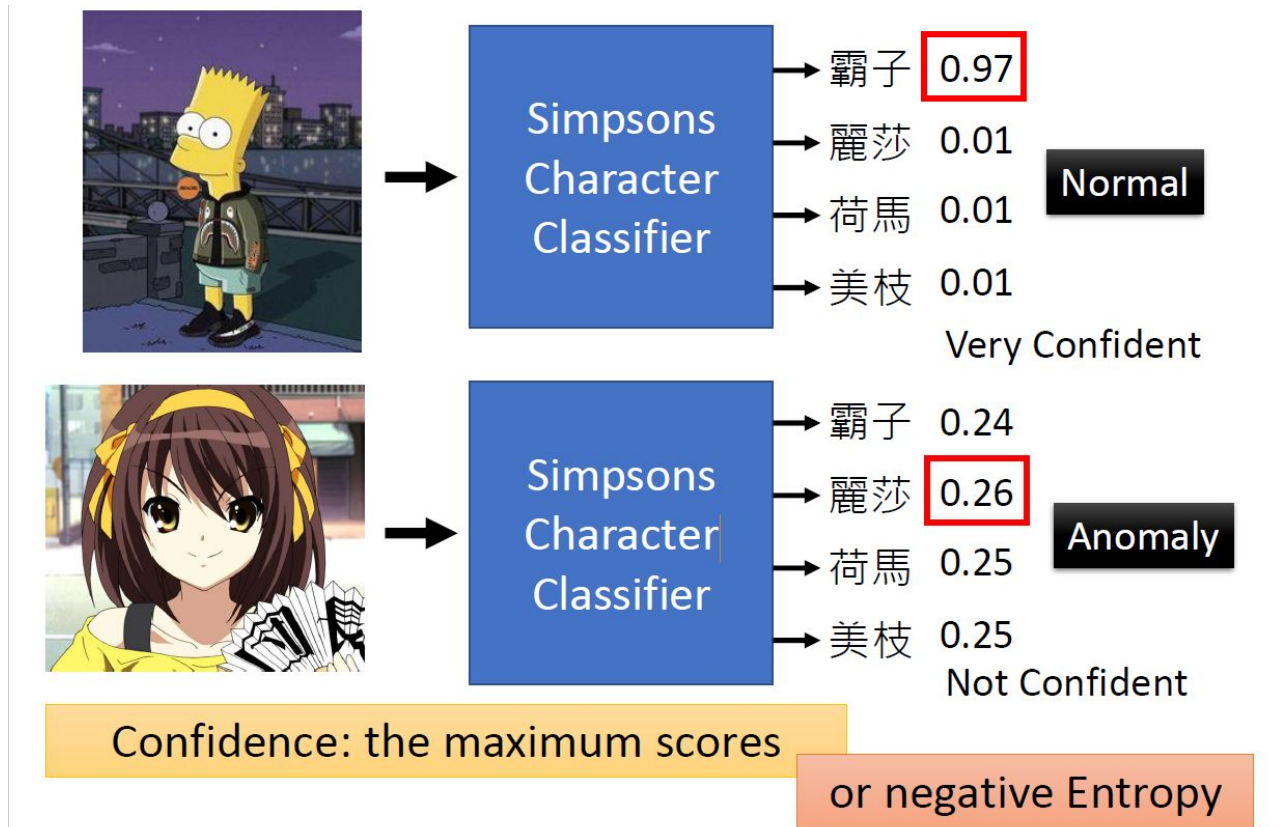
With Classifier



Anomaly Detection:


$$f(x) = \begin{cases} normal, & c(x) > \lambda \\ anomaly, & c(x) \leq \lambda \end{cases}$$

With Classifier







GMM (Gaussian Mixture Model)

1-dim Gaussian Mixtures



GMM (Gaussian Mixture Model)

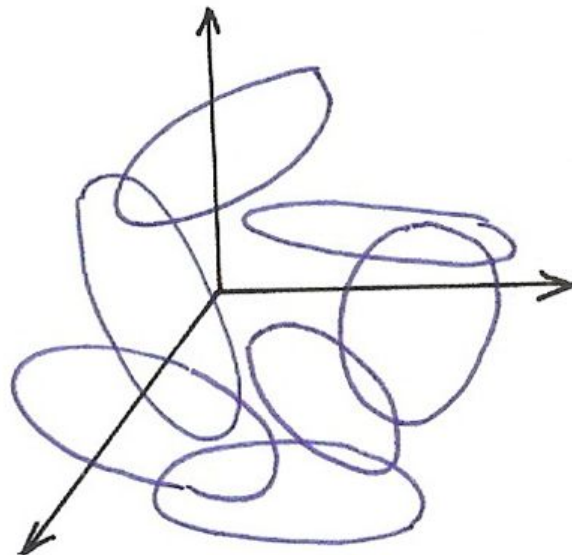
2-dim Gaussian



GMM (Gaussian Mixture Model)

N-dim Gaussian Mixtures

N-dim Gaussian Mixtures

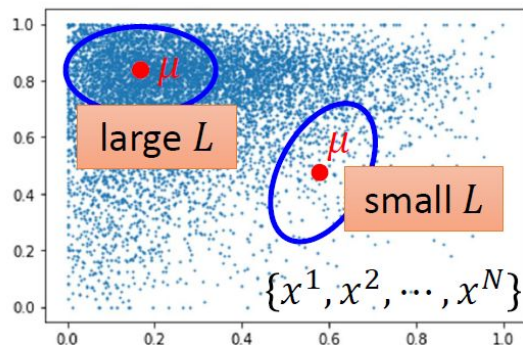


GMM (Gaussian Mixture Model)

$$f_{\mu, \Sigma}(x) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$$

Input: vector x , output: probability of sampling x

θ which determines the shape of the function are **mean** μ
and **covariance matrix** Σ



$$L(\theta) = f_{\theta}(x^1) f_{\theta}(x^2) \cdots f_{\theta}(x^N)$$

$$L(\mu, \Sigma) = f_{\mu, \Sigma}(x^1) f_{\mu, \Sigma}(x^2) \cdots f_{\mu, \Sigma}(x^N)$$

$$\theta^* = \arg \max_{\theta} L(\theta)$$

$$\mu^*, \Sigma^* = \arg \max_{\mu, \Sigma} L(\mu, \Sigma)$$

$$\mu^* = \frac{1}{N} \sum_{n=1}^N x^n = \begin{bmatrix} 0.29 \\ 0.73 \end{bmatrix}$$

$$\Sigma^* = \frac{1}{N} \sum_{n=1}^N (x - \mu^*)(x - \mu^*)^T = \begin{bmatrix} 0.04 & 0 \\ 0 & 0.03 \end{bmatrix}$$

GMM (Gaussian Mixture Model)

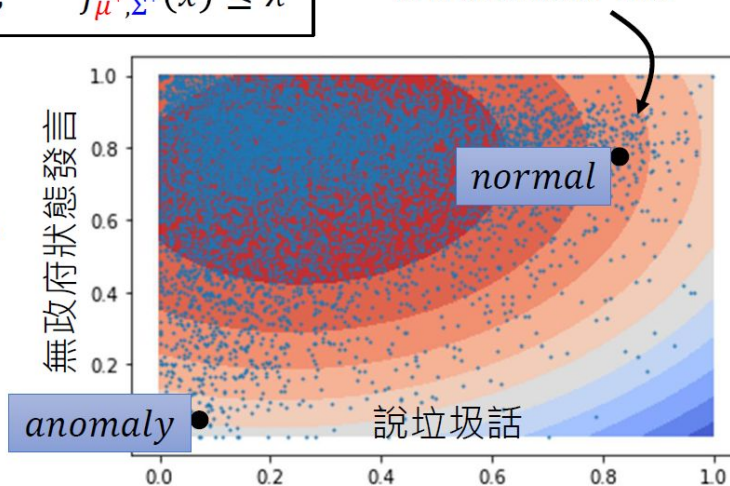
$$f_{\mu^*, \Sigma^*}(x) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma^*|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu^*)^T \Sigma^{*-1} (x - \mu^*) \right\}$$

$$\mu^* = \begin{bmatrix} 0.29 \\ 0.73 \end{bmatrix} \quad \Sigma^* = \begin{bmatrix} 0.04 & 0 \\ 0 & 0.03 \end{bmatrix}$$

$$f(x) = \begin{cases} \text{normal}, & f_{\mu^*, \Sigma^*}(x) > \lambda \\ \text{anomaly}, & f_{\mu^*, \Sigma^*}(x) \leq \lambda \end{cases}$$

λ is a contour line

The colors represents
the value of $f_{\mu^*, \Sigma^*}(x)$



GMM (Gaussian Mixture Model)

```
1  from keras.datasets import mnist
2  import cv2
3  import numpy as np
4
5  (x_train, y_train), (x_test, y_test) = mnist.load_data()
6  x_ok = x_train[y_train == 1] # 6742 筆
7  x_test = x_test[(y_test == 7) | (y_test == 1)] # 1135 筆 "1", 1028 筆 "7"
8  y_test = y_test[(y_test == 7) | (y_test == 1)]
9
10 def reshape_x(x):
11     new_x = np.empty((len(x), 56, 56))
12     for i, e in enumerate(x):
13         new_x[i] = cv2.resize(e, (56, 56))
14
15     new_x = np.expand_dims(new_x, axis=-1)
16     new_x = np.repeat(new_x, 3, axis=-1)
17     return new_x
18
19 x_ok = reshape_x(x_ok)
20 x_test = reshape_x(x_test)
```

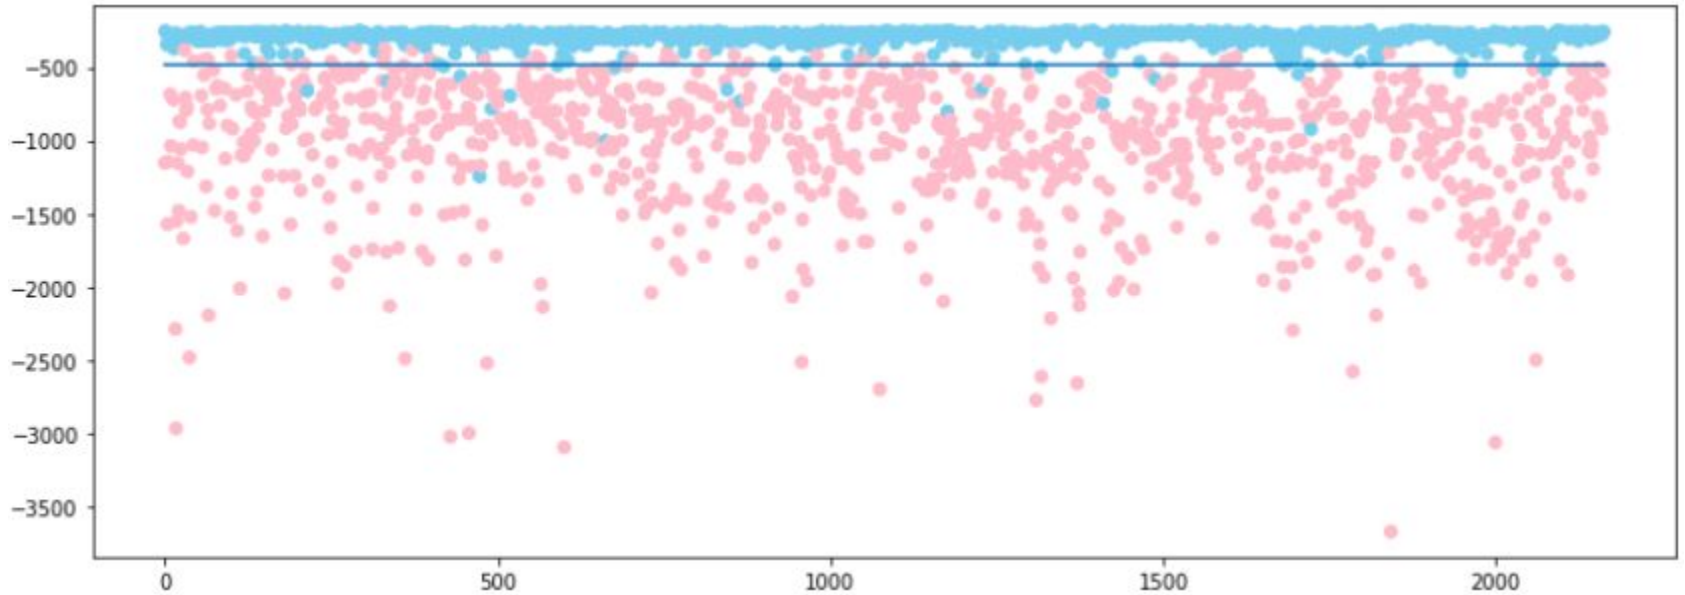
GMM (Gaussian Mixture Model)

```
features = model.predict(x_ok)
gmm.fit(features)
OKscore = gmm.score_samples(features)
thred = OKscore.mean() - 3 * OKscore.std()


test_features = model.predict(x_test)
score = gmm.score_samples(test_features)

print('normal accuracy: %.2f' % (len(score[(y_test == 1) & (score > thred)]) / 1135))
print('abnormal accuracy: %.2f' % (len(score[(y_test == 7) & (score < thred)]) / 1028))
      normal accuracy: 0.98
      abnormal accuracy: 0.96
```

GMM (Gaussian Mixture Model)



Auto-Encoder



Auto-Encoder

Examples:

<https://www.kaggle.com/tikedameu/anomaly-detection-with-autoencoder-pytorch>

<https://towardsdatascience.com/anomaly-detection-with-autoencoder-b4cdce4866a6>


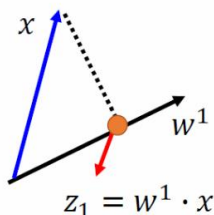
PCA

PCA

$$z = Wx$$

Reduce to 1-D:

$$z_1 = w^1 \cdot x$$



Project all the data points x onto w^1 , and obtain a set of z_1

We want the variance of z_1 as large as possible

$$\text{Var}(z_1) = \sum_{z_1} (z_1 - \bar{z}_1)^2 \quad \|w^1\|_2 = 1$$

PCA

$$z = Wx$$

Reduce to 1-D:

$$z_1 = w^1 \cdot x$$

$$z_2 = w^2 \cdot x$$

$$W = \begin{bmatrix} (w^1)^T \\ (w^2)^T \\ \vdots \end{bmatrix}$$

Orthogonal matrix

Project all the data points x onto w^1 , and obtain a set of z_1

We want the variance of z_1 as large as possible

$$\text{Var}(z_1) = \sum_{z_1} (z_1 - \bar{z}_1)^2 \quad \|w^1\|_2 = 1$$

We want the variance of z_2 as large as possible

$$\text{Var}(z_2) = \sum_{z_2} (z_2 - \bar{z}_2)^2 \quad \|w^2\|_2 = 1$$

$$w^1 \cdot w^2 = 0$$

PCA

```
# 30 principal components
from sklearn.decomposition import PCA

n_components = 30
whiten = False
random_state = 2018

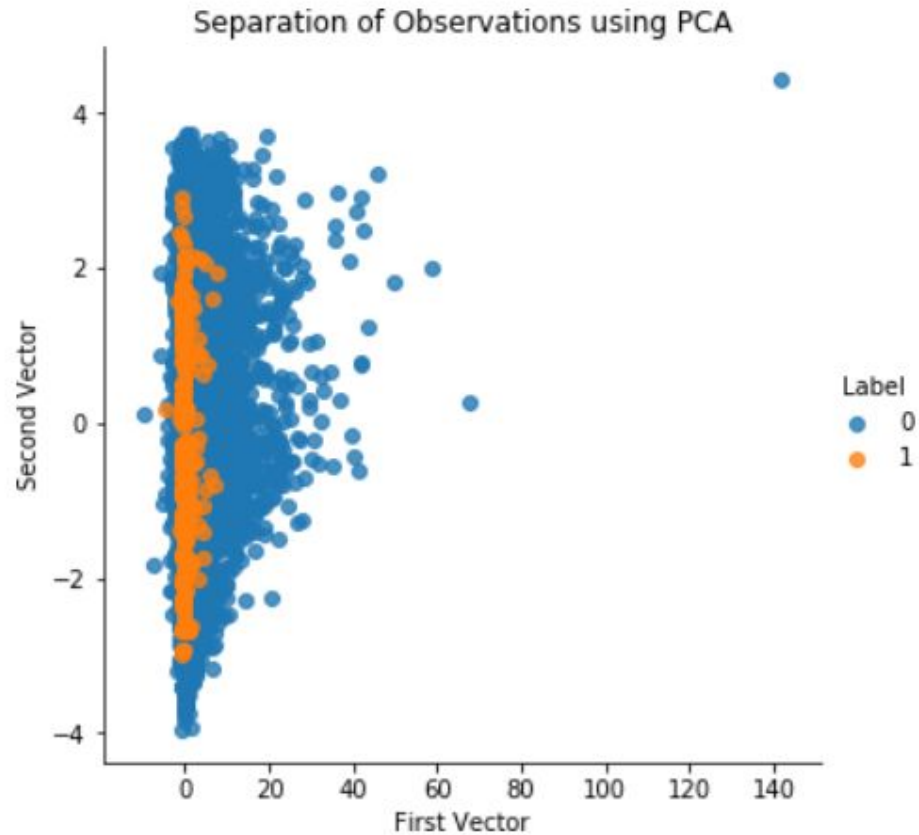
pca = PCA(n_components=n_components, whiten=whiten, \
          random_state=random_state)

X_train_PCA = pca.fit_transform(X_train)
X_train_PCA = pd.DataFrame(data=X_train_PCA, index=X_train.index)

X_train_PCA_inverse = pca.inverse_transform(X_train_PCA)
X_train_PCA_inverse = pd.DataFrame(data=X_train_PCA_inverse, \
                                   index=X_train.index)

scatterPlot(X_train_PCA, y_train, "PCA")
```

PCA

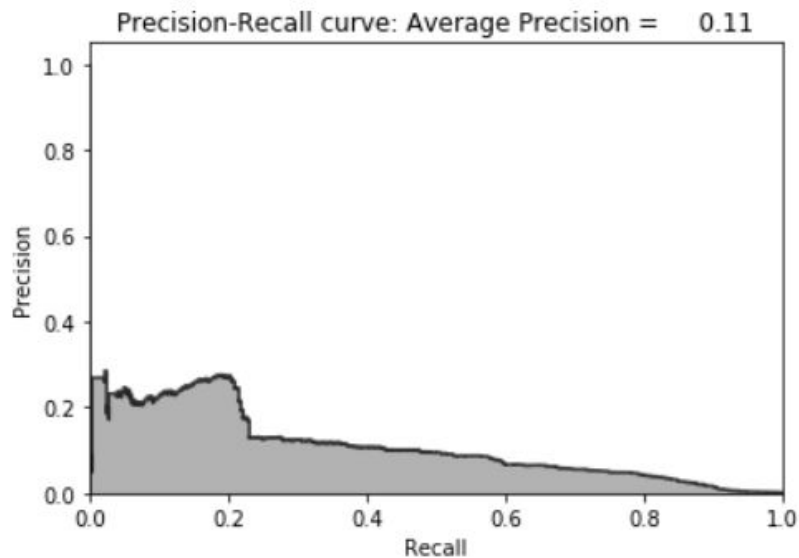


PCA

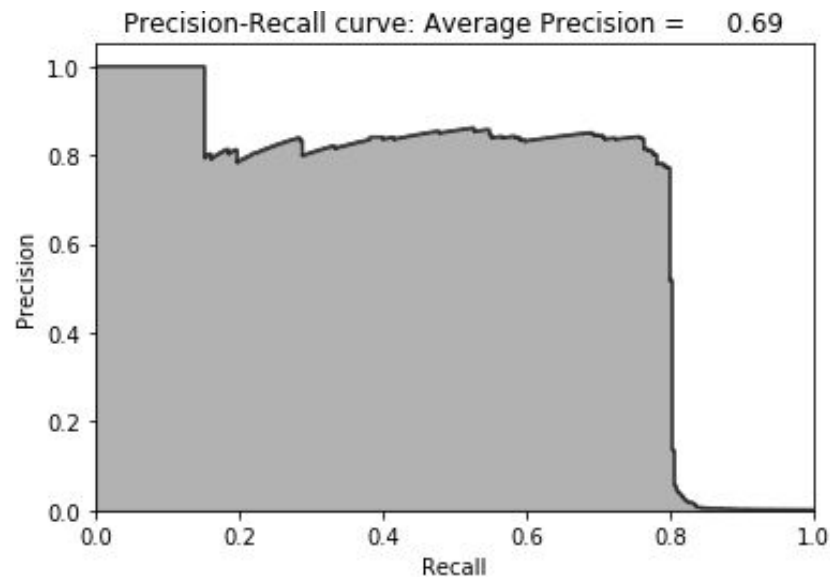
```
def anomalyScores(originalDF, reducedDF):  
    loss = np.sum((np.array(originalDF) - np.array(reducedDF))**2, axis=1)  
    loss = pd.Series(data=loss, index=originalDF.index)  
    loss = (loss - np.min(loss)) / (np.max(loss) - np.min(loss))  
    return loss
```

```
anomalyScoresPCA = anomalyScores(X_train, X_train_PCA_inverse)  
preds = plotResults(y_train, anomalyScoresPCA, True)
```

PCA



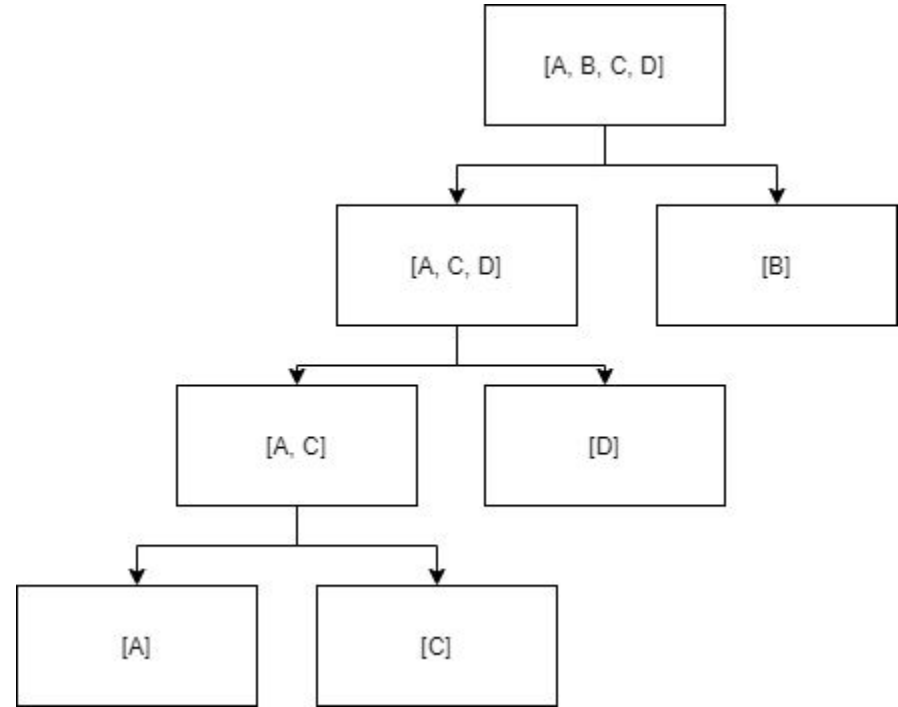
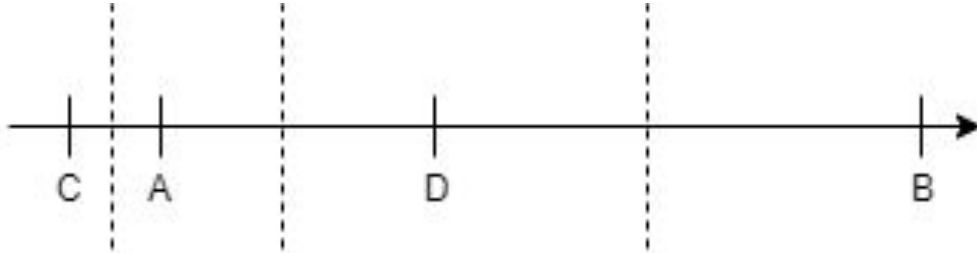
n_components = 30



n_components = 27

Isolation Forest

- Use tree like structure to split data



Isolation Forest

```
# importing libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from pylab import savefig
from sklearn.ensemble import IsolationForest

# Generating data ----

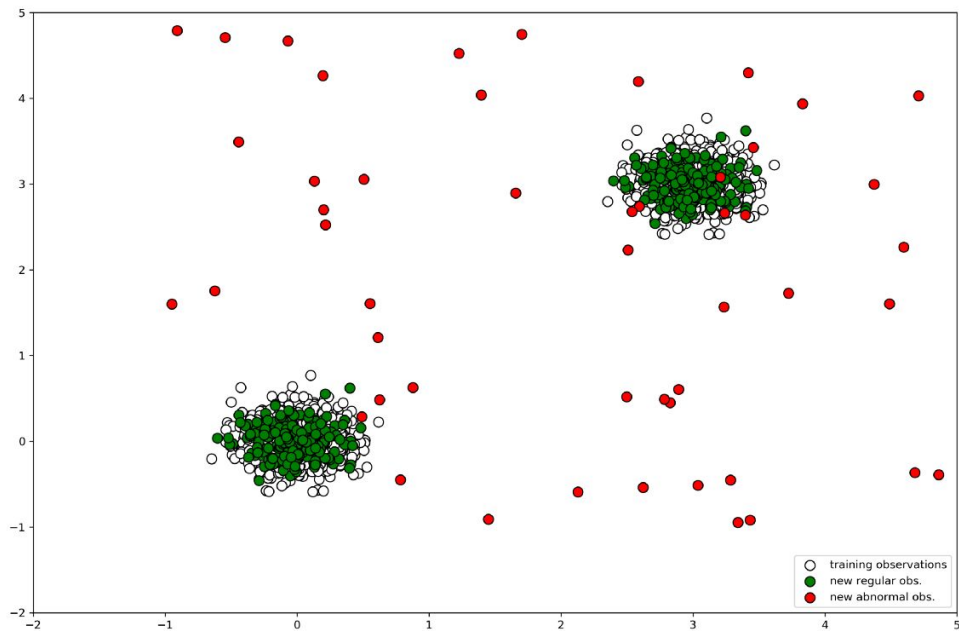
rng = np.random.RandomState(42)

# Generating training data
X_train = 0.2 * rng.randn(1000, 2)
X_train = np.r_[X_train + 3, X_train]
X_train = pd.DataFrame(X_train, columns = ['x1', 'x2'])

# Generating new, 'normal' observation
X_test = 0.2 * rng.randn(200, 2)
X_test = np.r_[X_test + 3, X_test]
X_test = pd.DataFrame(X_test, columns = ['x1', 'x2'])

# Generating outliers
X_outliers = rng.uniform(low=-1, high=5, size=(50, 2))
X_outliers = pd.DataFrame(X_outliers, columns = ['x1', 'x2'])
```


Isolation Forest



Isolation Forest

```
# Isolation Forest ----  
  
# training the model  
clf = IsolationForest(max_samples=100, random_state=rng)  
clf.fit(X_train)  
  
# predictions  
y_pred_train = clf.predict(X_train)  
y_pred_test = clf.predict(X_test)  
y_pred_outliers = clf.predict(X_outliers)
```

Isolation Forest

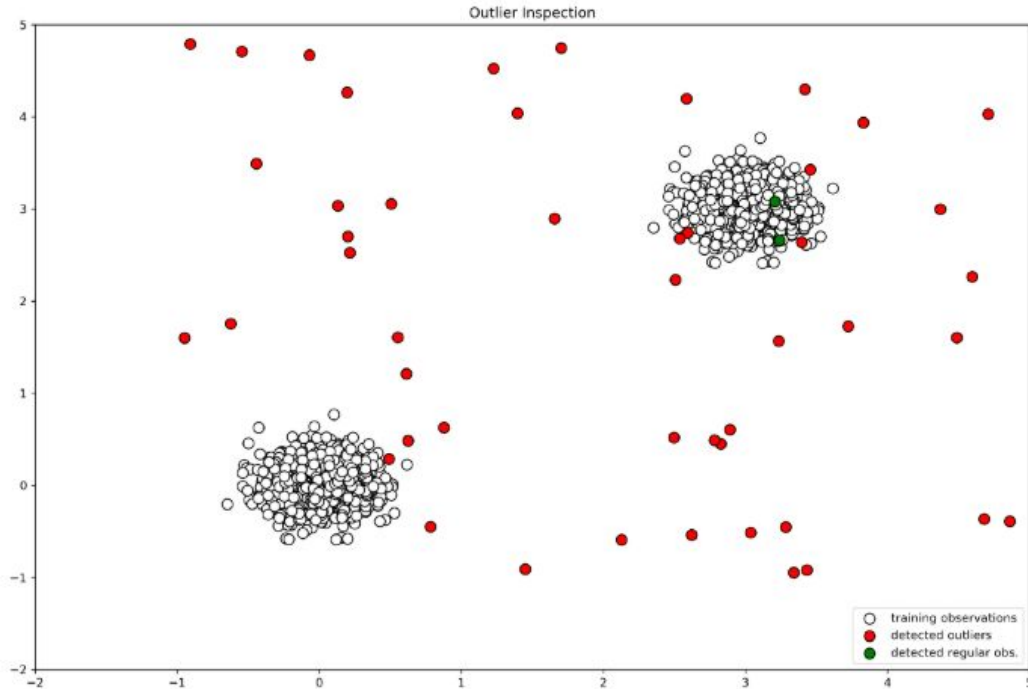


Figure 3 Inspecting outlier classification

<https://towardsdatascience.com/outlier-detection-with-isolation-forest-3d190448d45e>

Summary of classic method

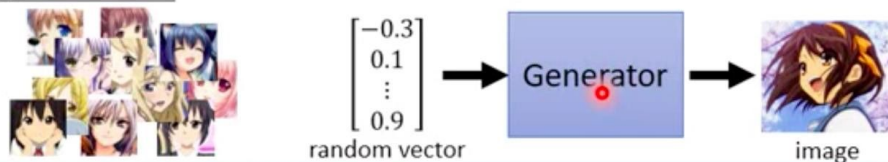
- With Classifier
- GMM (Gaussian Mixture Model)
- Auto-Encoder
- PCA
- Isolation Forest

Anomaly detection on image

Typical GANs

Three Categories of GAN

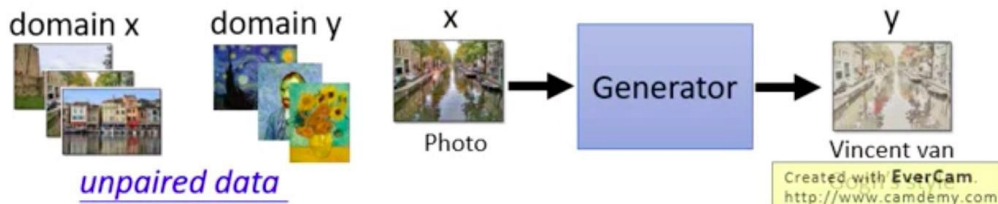
1. Typical GAN



2. Conditional GAN



3. Unsupervised Conditional GAN



Typical GANs

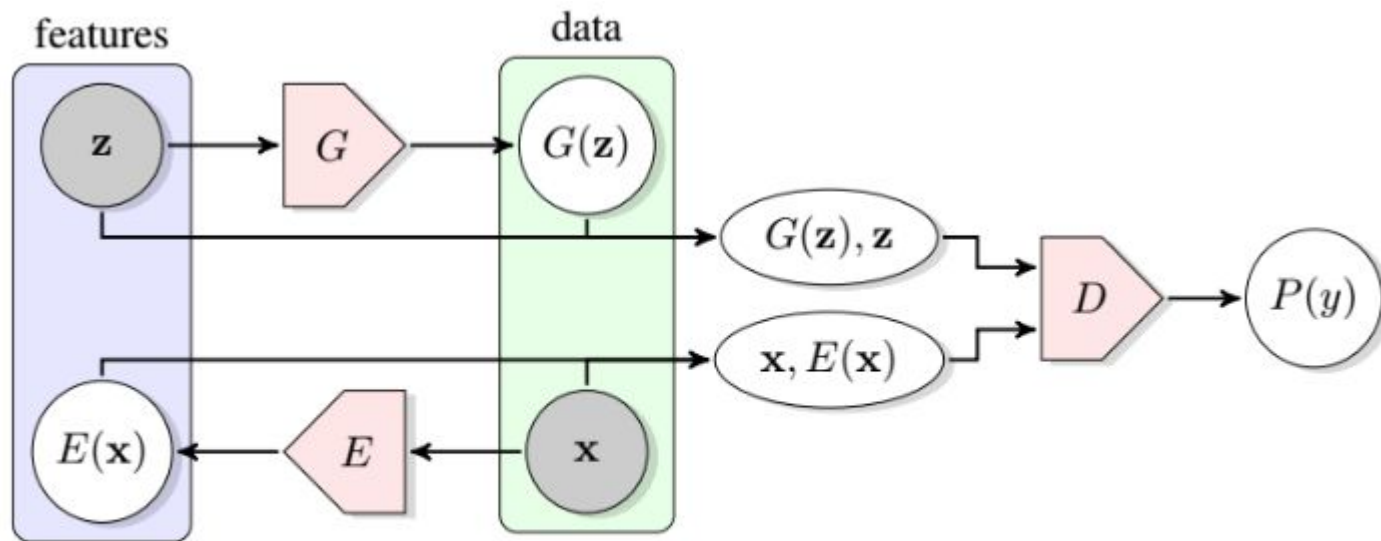



Figure 1. The structure of BiGAN proposed in (Donahue et al., 2016).

Example

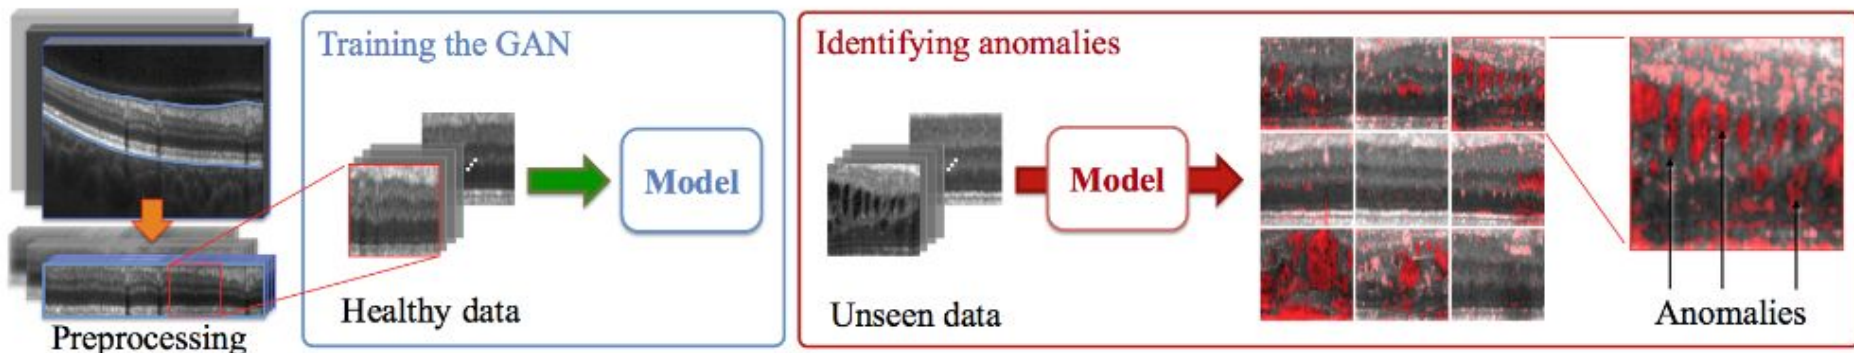


(a) Normal Data (X-ray Scans)

(b) Normal + Abnormal Data (X-ray Scans)

AnoGAN

- Train a standard GAN only on positive samples



AnoGAN

- Anomaly Score

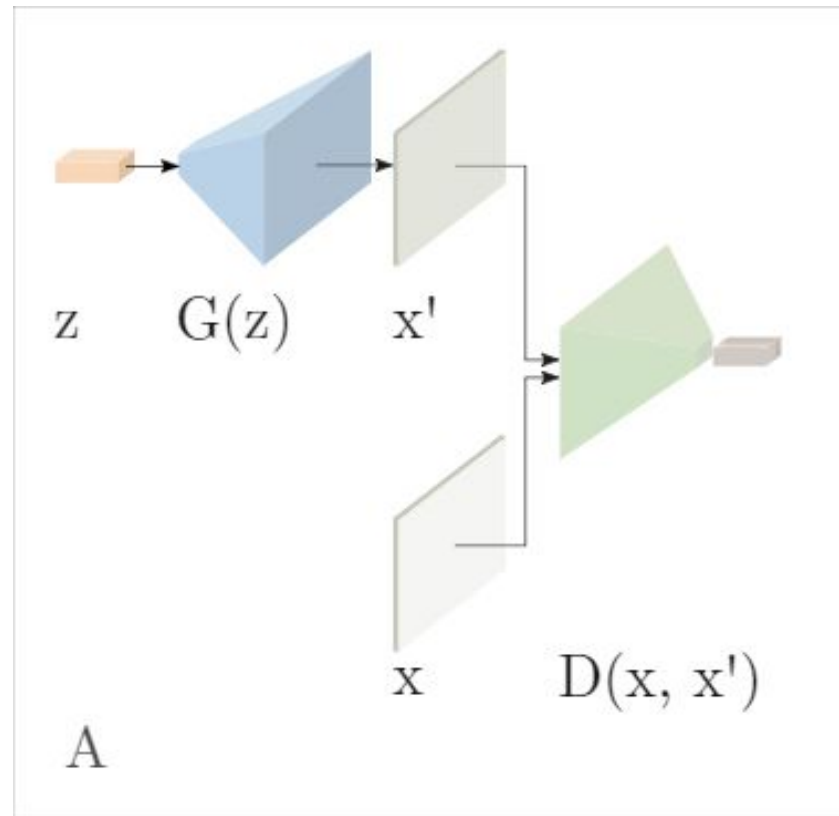
$$A(\mathbf{x}) = \mathcal{L}(\mathbf{z}_\Gamma)$$

- for $\gamma = 1, 2, \dots, \Gamma$ find proper \mathbf{z}

$$\mathcal{L}_R(\mathbf{z}_\gamma) = \|\mathbf{x} - G(\mathbf{z}_\gamma)\|_1$$

$$\mathcal{L}_D(\mathbf{z}_\gamma) = \|\mathbf{f}(\mathbf{x}) - \mathbf{f}(G(\mathbf{z}_\gamma))\|_1$$

$$\mathcal{L}(\mathbf{z}_\gamma) = (1 - \lambda) \cdot \mathcal{L}_R(\mathbf{z}_\gamma) + \gamma \cdot \mathcal{L}_D(\mathbf{z}_\gamma)$$



AnoGAN

- Pros

- Showed that GANs can be used for anomaly detection
- Introduced a new mapping scheme from latent space to input data space.
- Used the same mapping scheme to define an anomaly score.

- Cons

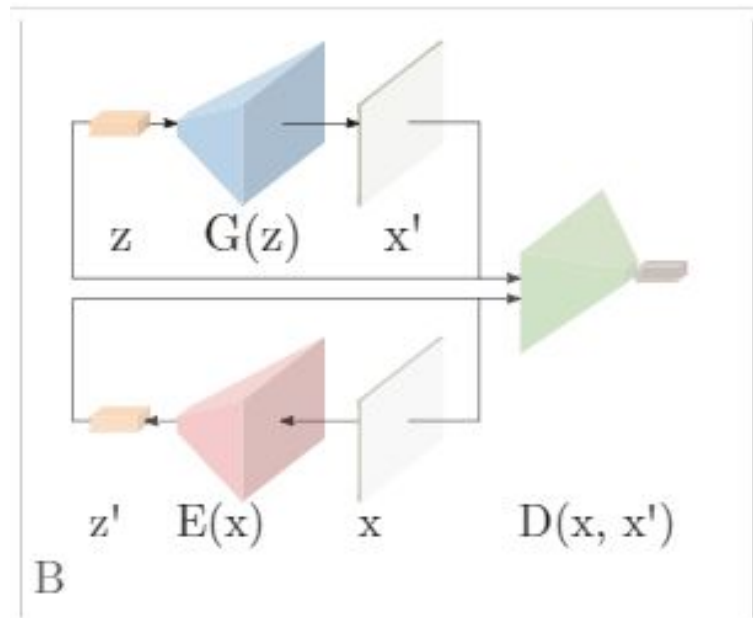
- Requires Γ optimization steps for every new input: bad test-time performance
- The GAN objective has not been modified to take into account the need for the inverse mapping learning.
- The anomaly score is difficult to interpret, not being in the probability range.

EGBAD (Efficient GAN-Based Anomaly Detection)

- Train a Bi-GAN only on positive samples

$$\min_{G,E} \max_D V(D, G, E) =$$

$$\mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\mathbb{E}_{\mathbf{z} \sim p_{E(\mathbf{z}|\mathbf{x})}} [\log D(\mathbf{x}, \mathbf{z})]] +$$
$$\mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\mathbb{E}_{\mathbf{x} \sim p_G(\mathbf{x}|\mathbf{z})} [\log(1 - D(\mathbf{x}, \mathbf{z}))]].$$



EGBAD (Efficient GAN-Based Anomaly Detection)

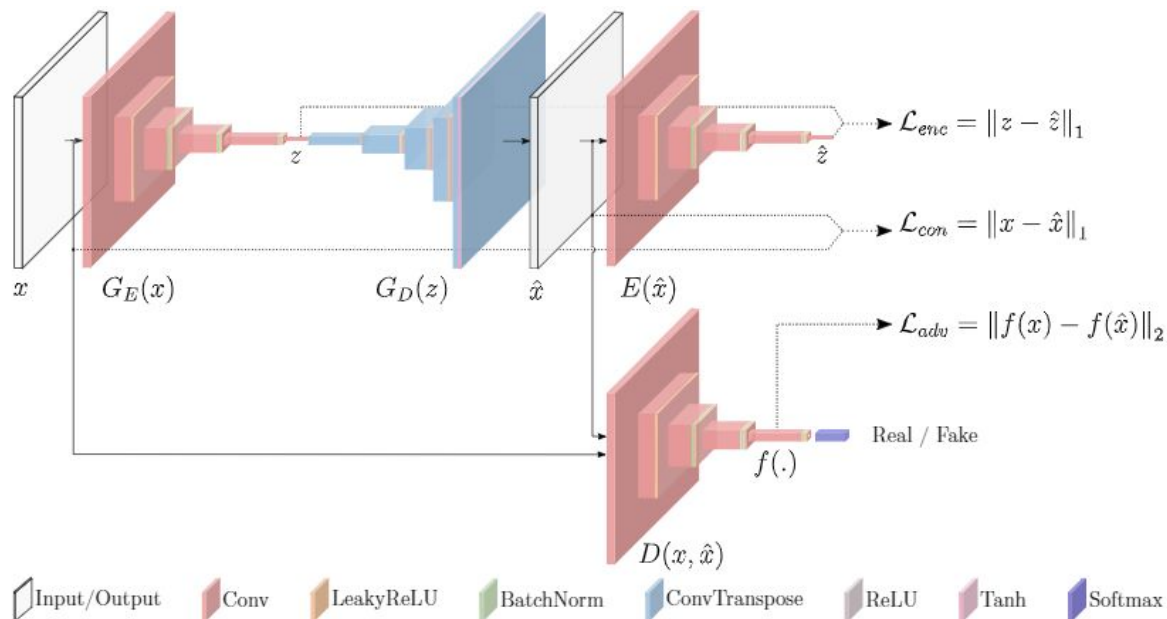
- Pros
 - The Encoder E can learn how to encode image while adversarial training.
 - Therefore, it can bypass Γ optimization steps of AnoGAN while calculating anomaly score.

GANomaly

- Generator is composed of encoder G_E , decoder G_D , and encoder E
- Trained on only normal data.

- Anomaly score

$$\mathcal{A}(\mathbf{x}) = \|G_E(\mathbf{x}) - E(G(\mathbf{x}))\|_2$$



GANomaly

- Pros

- An encoder is learned during the training process, so it can bypass the Γ optimization.
- Using an autoencoder like architecture (no use of noise prior) makes the entire learning process faster.
- The contextual loss can be used to localize the anomaly.

- Cons

- Defines a new anomaly score.
- It allows to detect anomalies both in the image space and in the latent space, but the results couldn't match:
 - a higher anomaly score, that's computed only in the latent space, can be associated with a generated sample with a low contextual loss value and thus very similar to the input - and vice versa.


Evaluation Metric

- TPR (True Positive Rate)

$$\frac{TP}{TP + FN}$$

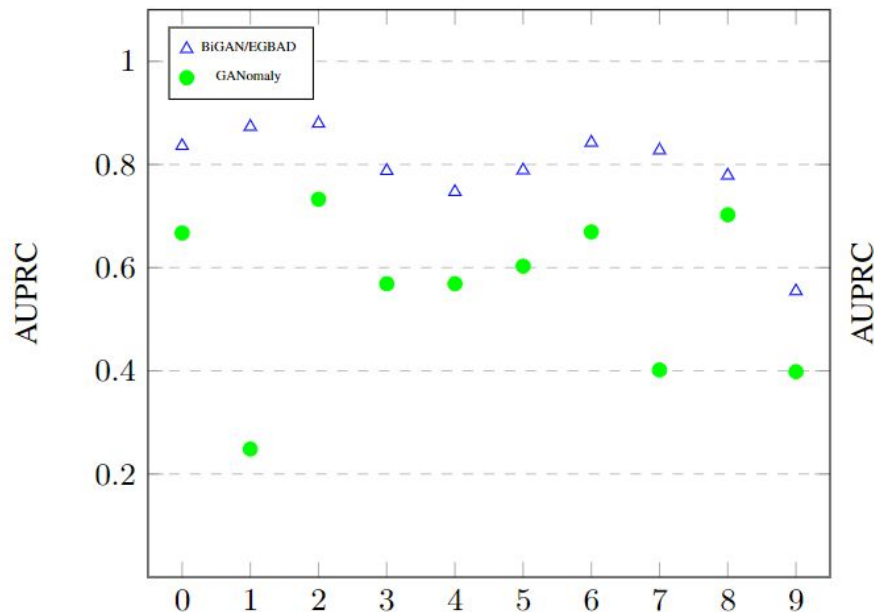
- FPR (False Positive Rate)

$$\frac{FP}{TN + FP}$$

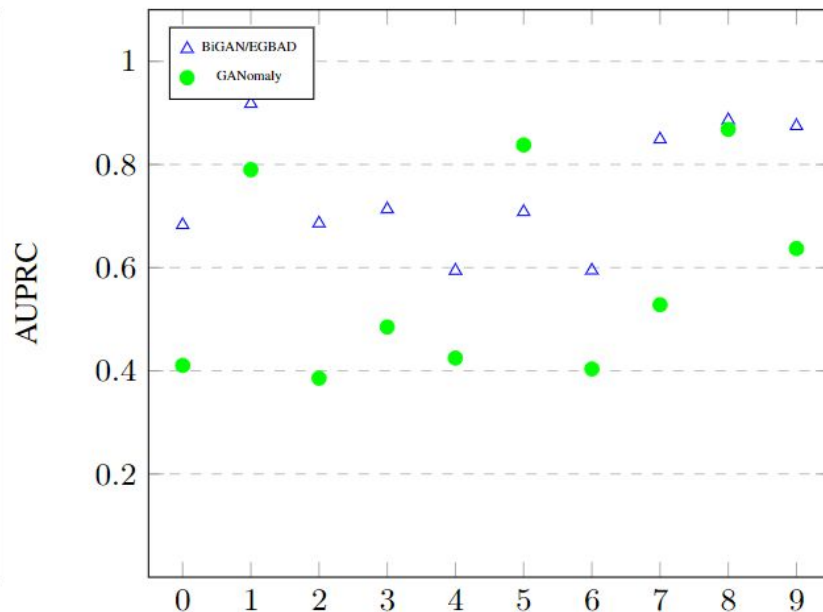


Comparison

BiGAN/EGBAD and GANomaly performance comparison on MNIST and Fashion-MNIST datasets


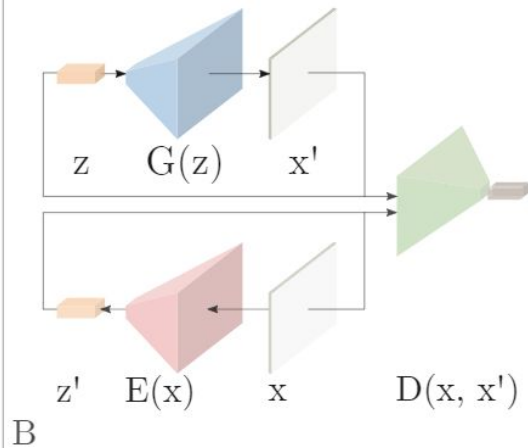
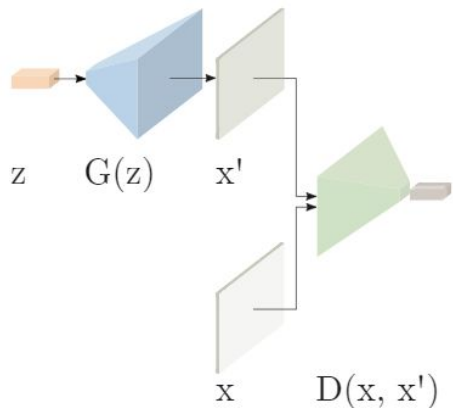


(a) MNIST anomalous classes



(b) Fashion-MNIST anomalous classes

Summary of GANs



AnoGAN

EGBAD

GANomaly

Reference

<https://arxiv.org/pdf/1711.09325.pdf>

<https://arxiv.org/pdf/1809.10816.pdf>

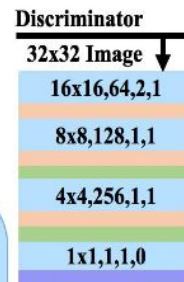
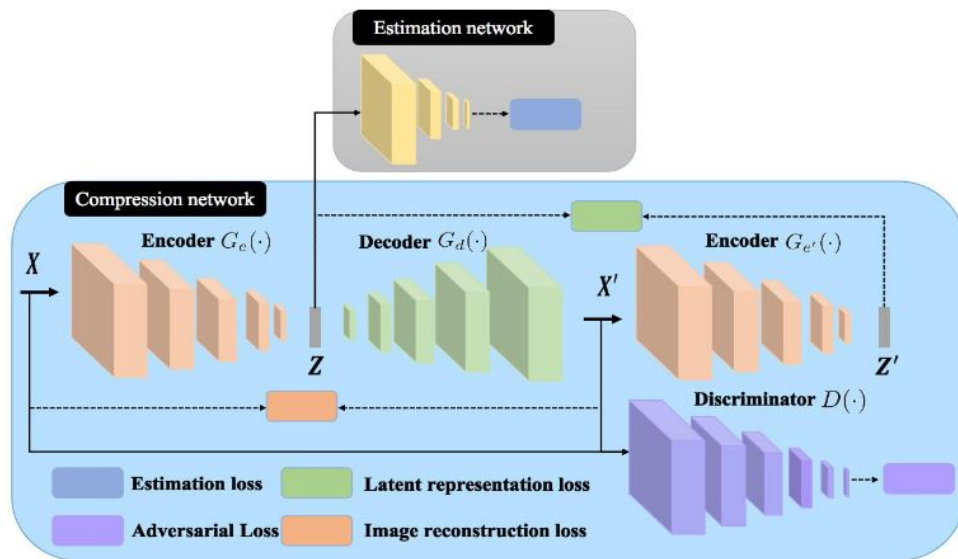
<https://arxiv.org/pdf/1812.02288.pdf>

<https://arxiv.org/pdf/1901.08954.pdf>




<https://arxiv.org/pdf/1905.13147.pdf>

Anomaly detection on Audio

GMGAN (Gaussian Mixture GAN)



Fake / Real



GMGAN (Gaussian Mixture GAN)

- Adversarial loss

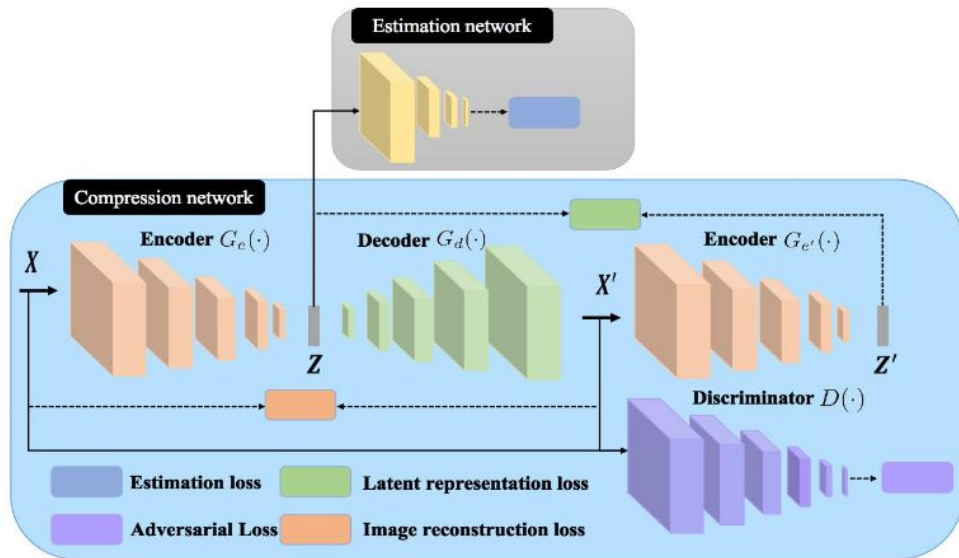
$$\mathcal{L}_{adv} = \min_G \max_D (E_{\mathbf{x} \sim p_{\mathbf{x}}} [\log(D(\mathbf{x}))] + E_{\mathbf{x} \sim p_{\mathbf{x}}} [\log(1 - D(G(\mathbf{x})))]).$$

- Image reconstruction loss

$$\mathcal{L}_{irec} = \mathbb{E}_{x \sim p_{\mathbf{x}}} \|x - G(x)\|_1$$

- Latent representation loss

$$\mathcal{L}_{zrec} = \mathbb{E}_{x \sim p_{\mathbf{X}}} \|G_e(x) - G_{e'}(x')\|_2$$

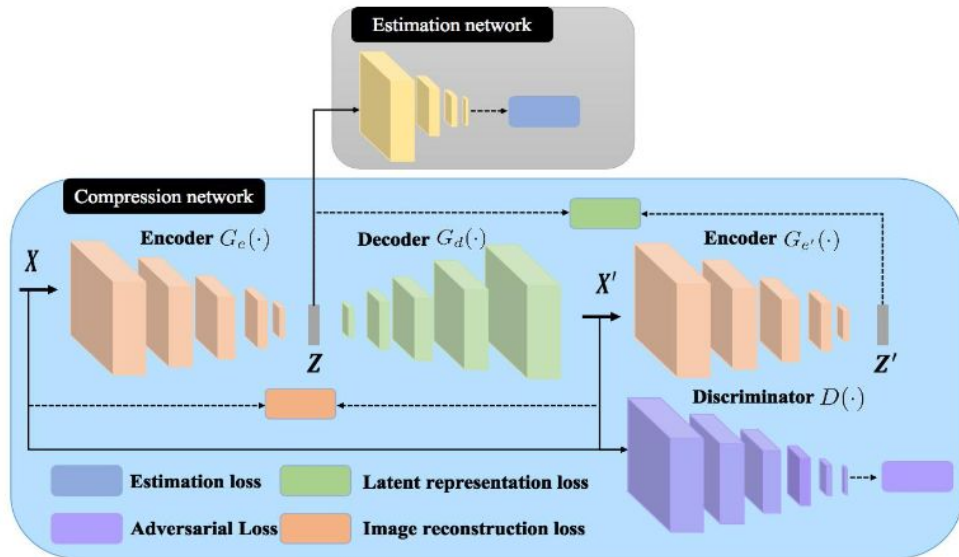


GMGAN (Gaussian Mixture GAN)

- Estimation loss
- $\hat{\gamma} = \text{softmax}(MLN(\mathbf{z}; \theta_m))$
- $\hat{\gamma}$ is a K-dimensional vector and $\hat{\gamma}_k$ denotes the input belongs to the k^{th} distribution.
- Calculate the component of k^{th} mixture.

$$\hat{\alpha}_k = \frac{1}{n} \sum_{i=1}^n \hat{\gamma}_{ik}; \hat{u}_k = \frac{\sum_{i=1}^n \hat{\gamma}_{ik} z_i}{\sum_{i=1}^n \hat{\gamma}_{ik}},$$

$$\hat{\Sigma}_k = \frac{\sum_{i=1}^n \hat{\gamma}_{ik} (z_i - \hat{u}_k)(z_i - \hat{u}_k)^T}{\sum_{i=1}^n \hat{\gamma}_{ik}},$$

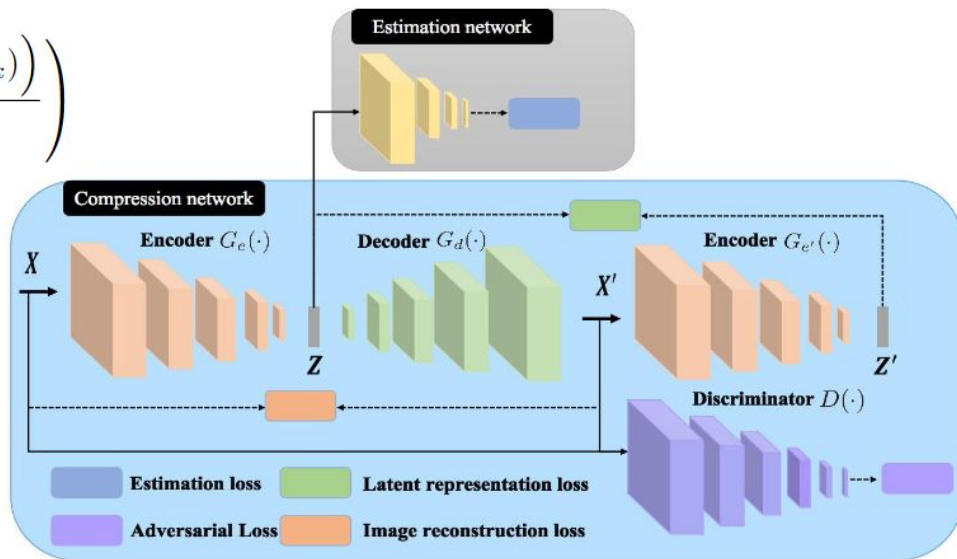


GMGAN (Gaussian Mixture GAN)

- Estimation loss

$$E(\mathbf{z}) = -\log \left(\sum_{k=1}^K \hat{\alpha}_k \frac{\exp \left(-\frac{1}{2} (\mathbf{z} - \hat{\mu}_k)^T \hat{\Sigma}_k^{-1} (\mathbf{z} - \hat{\mu}_k) \right)}{\sqrt{|2\pi \Sigma_k|}} \right)$$

$$\mathcal{L}_{es} = \lambda_1 \sum_{i=1}^N E(z_i) + \lambda_2 \sum_{k=1}^K \sum_{j=1}^d \frac{1}{\hat{\Sigma}_{jj}^k},$$



GMGAN (Gaussian Mixture GAN)




Fig. 2. Spectrogram of one second sequence containing office background audio and glass break audio.

GMGAN (Gaussian Mixture GAN)

Scene	CAE [9]	WaveNet [19]	Proposed method
beach	0.69	0.72	0.80
bus	0.79	0.83	0.89
cafe/restaurant	0.69	0.76	0.76
car	0.79	0.82	0.93
city center	0.75	0.82	0.83
forest path	0.65	0.72	0.77
grocery store	0.71	0.77	0.83
home	0.69	0.69	0.69
library	0.59	0.67	0.85
metro station	0.74	0.79	0.81
office	0.78	0.78	0.80
park	0.70	0.80	0.89
residential area	0.73	0.78	0.78
train	0.82	0.84	0.92
tram	0.80	0.87	0.94

Outline

- What is Anomaly Detection
- Classic Method
 - With Classifier
 - GMM (Gaussian Mixture Model)
 - Auto-Encoder
 - PCA
 - Isolation Forest
 - Summary
- Anomaly Detection with GAN
 - AnoGAN
 - EGBAD
 - GANomaly
 - Summary
- Anomaly Detection on Audio
 - GMGAN

More Reference

Anomaly detection on self-driving

<http://taoxie.cs.illinois.edu/publications/tii17-safedrive.pdf>

<https://arxiv.org/pdf/2004.09496.pdf>

<https://arxiv.org/pdf/2004.12581.pdf>